Big Data and Hadoop

Venkatesh Vinayakarao
venkateshv@cmi.ac.in
http://vvtesh.co.in

Chennai Mathematical Institute

Data is the new oil. - Clive Humby, 2006.

What Comes Next?

byte

kilobyte

megabyte

gigabyte

??

???

????

?????

Sizes

Name	Size
Byte	8 bits
Kilobyte	1024 bytes
Megabyte	1024 kilobytes
Gigabyte	1024 megabytes
Terabyte	1024 gigabytes
Petabyte	1024 terabytes
Exabyte	1024 petabytes
Zettabyte	1024 exabytes
Yottabyte	1024 zettabytes

Data Growth

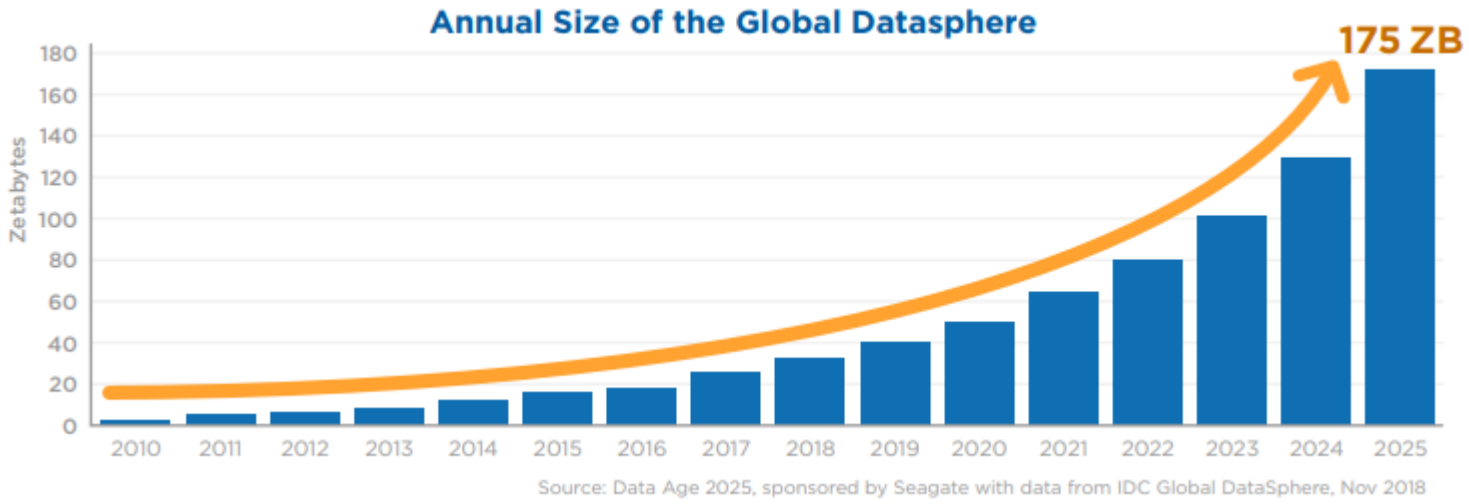

Mankind's quest to digitize the world!
33 ZB (2018) → 175 ZB (2025)
size of global datasphere*

*Source: https://www.seagate.com/files/www-content/our-story/trends/files/idc-seagate-dataage-whitepaper.pdf

Evolution of Data and Computers

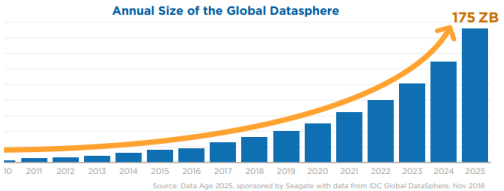

Von Neumann Arch

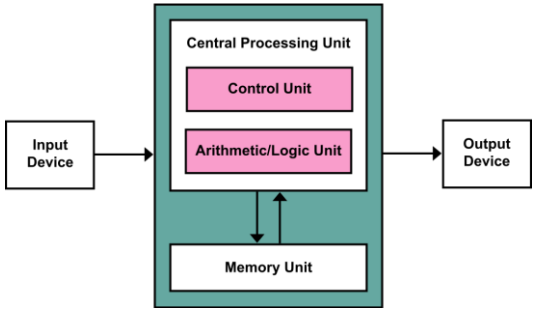

Challenges

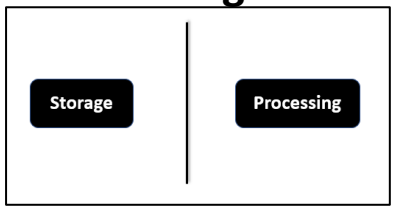

Data Storage

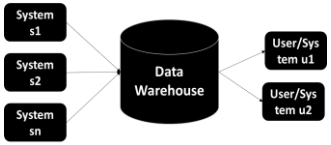

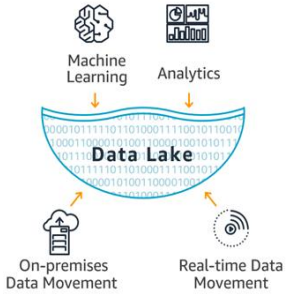

Amazon S3

STaaS

Recap

Data Storage

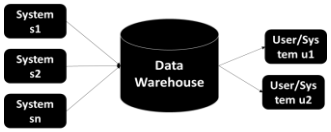

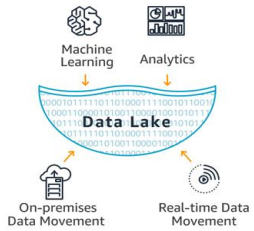

Amazon S3

STaaS

Data Processing

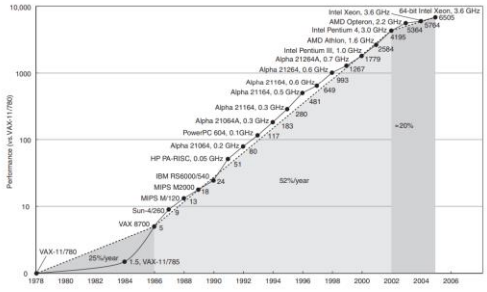

CPU Performance

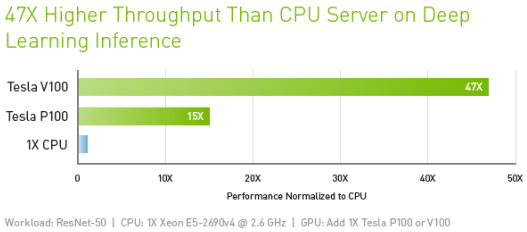

GPU Performance

SuperComputers

Cloud Computing

Two kinds of Big Data Opportunities

Storage

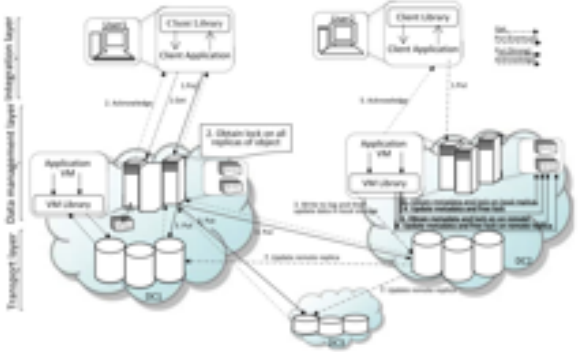

Processing

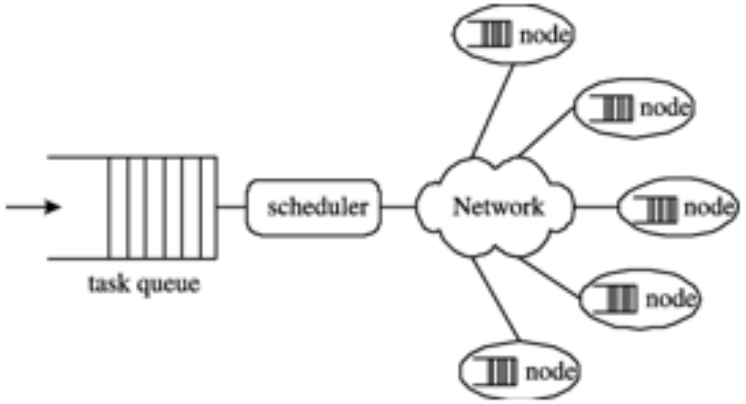

So, we have the cloud. But, how to store and retrieve data? How to process jobs?

Role of File Systems

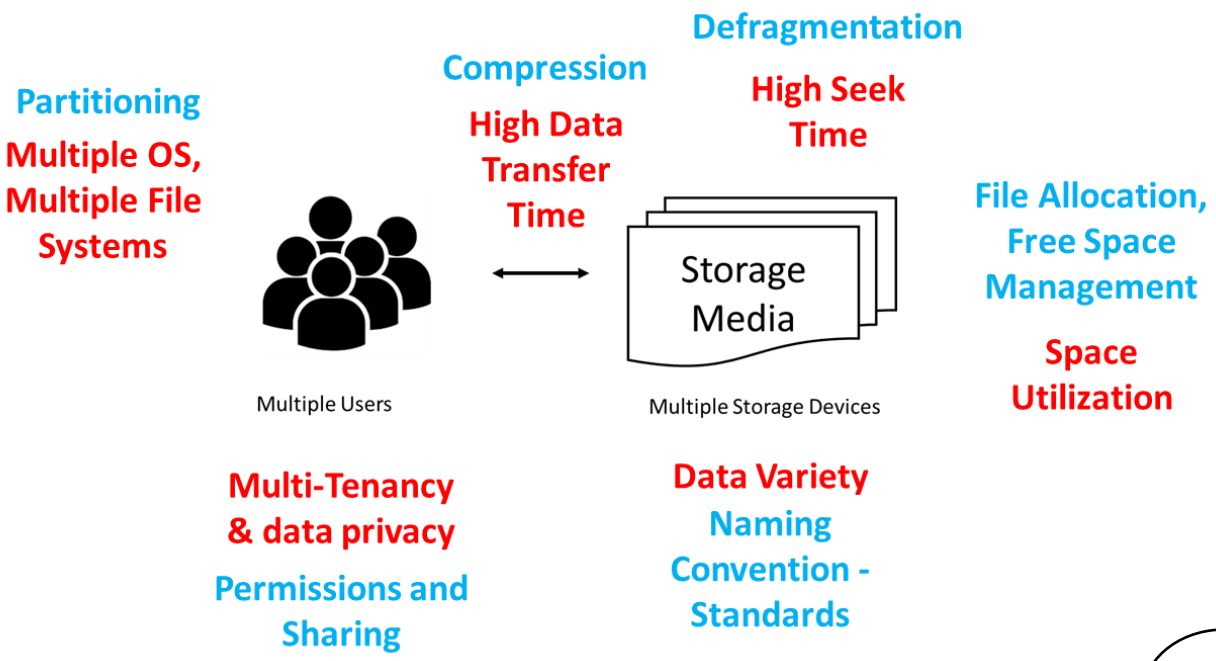

File systems are key to handling data.

Variety of FS exist
NTFS, FAT, DOS, CDFS, NFS, ...

Distributed Systems

Not designed for co-ordination jobs.

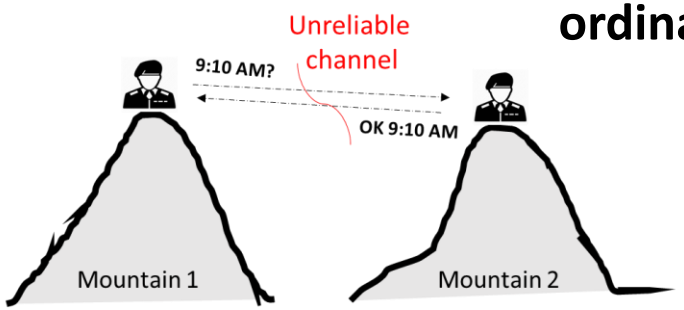

General's Paradox

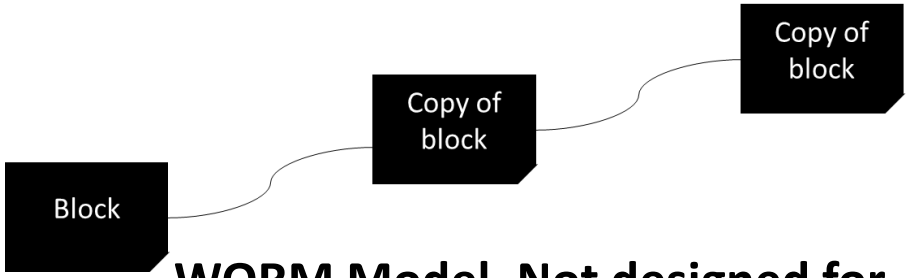

WORM Model. Not designed for write-many (interactive) jobs.

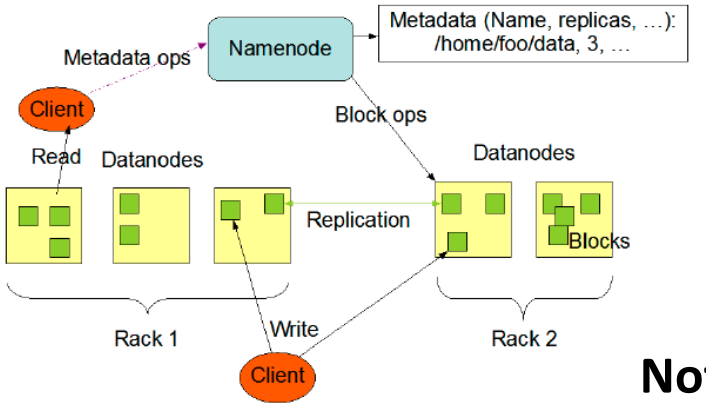

Not designed for small files.

Hadoop and Map Reduce

When not to use Hadoop?

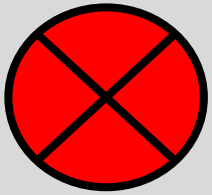

No Interactive Jobs
No Jobs Requiring Co-ordination
No Small Files

Hadoop Architecture

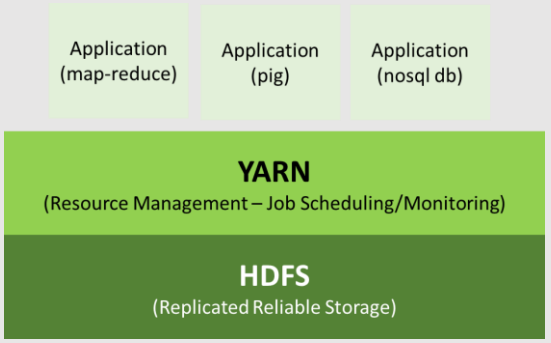

Map-reduce Model

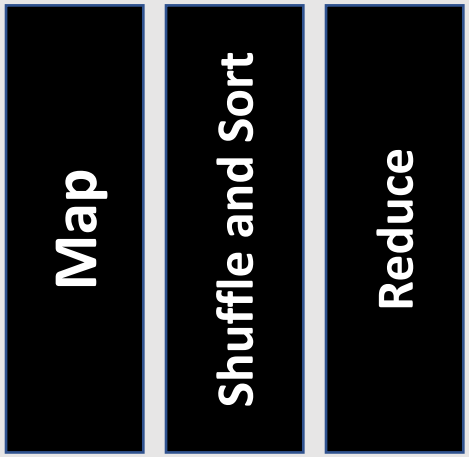

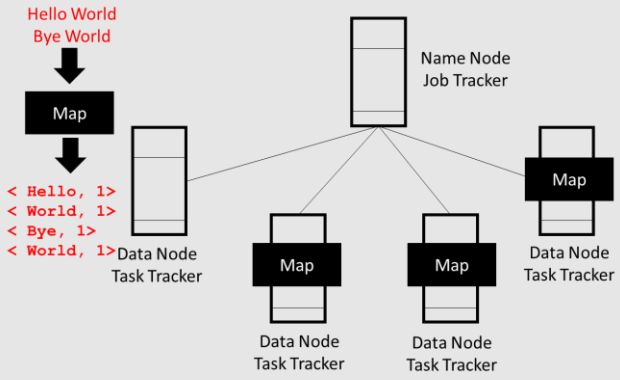

Map-Reduce Patterns

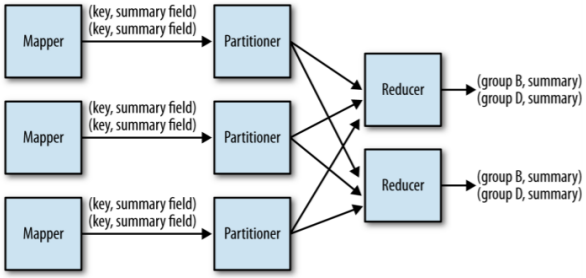

Summarization

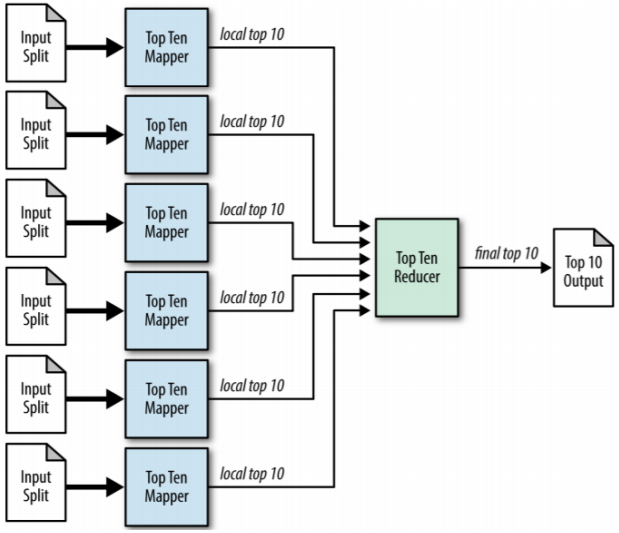

Top 10

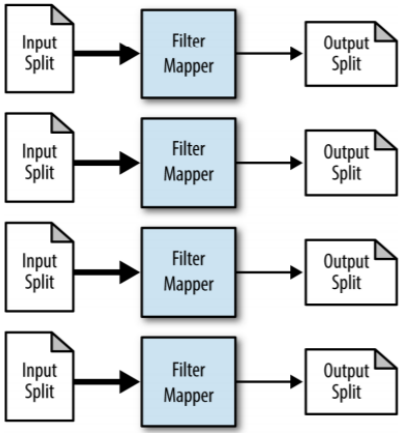

Filtering

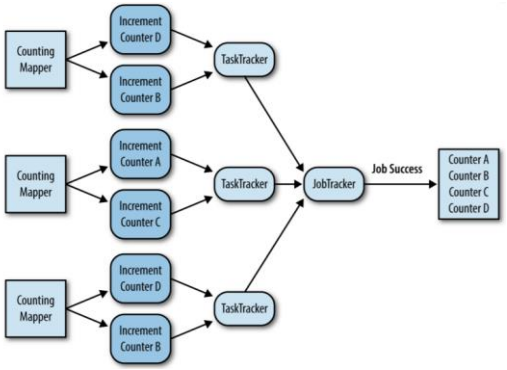

Counting

NoSQL

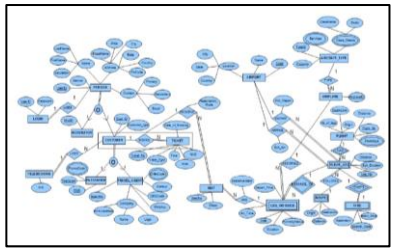

Schema-based Relational Model - maintenance problems

Impedance Mismatch

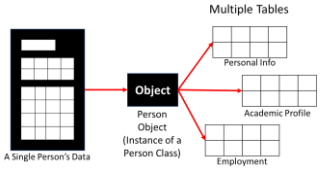

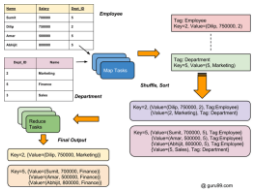

Scale-up Challenges

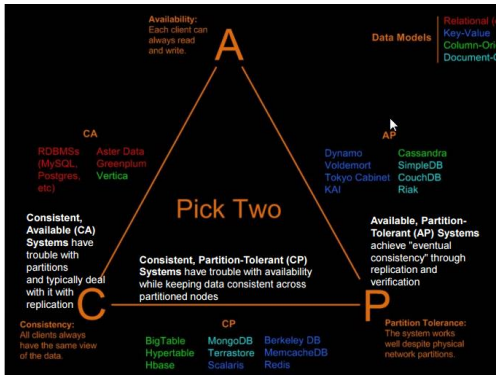

CAP Theorem

Types of NoSQL datastores

Key-Valuecv

```
redis> GET nonexisting
(nil)
redis> SET mykey "Hello"
"OK"
redis> GET mykey
"Hello"
redis>
```

redis

Doc-based

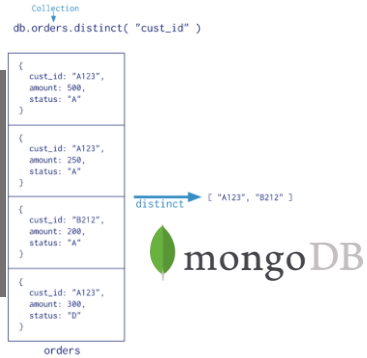

mongoDB

Columnar DB

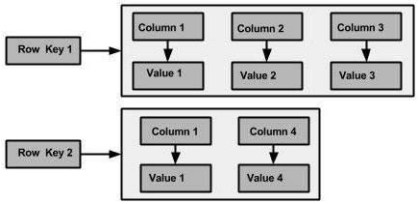

Apache CASSANDRA

Graph DB

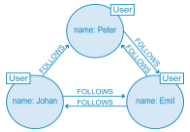

neo4j

Web Services

Interoperability

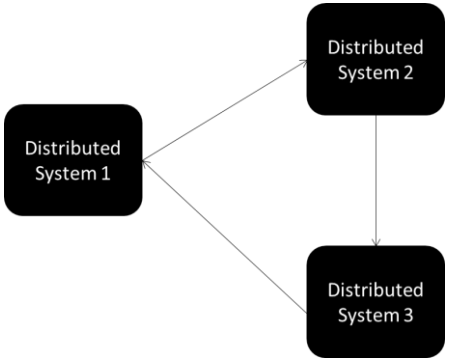

CORBA

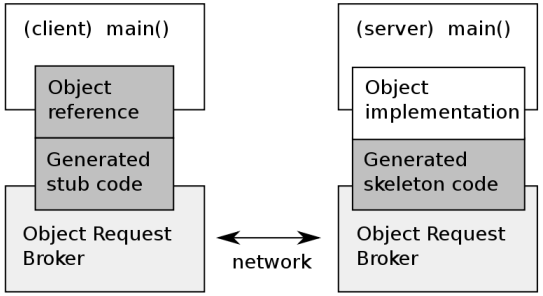

RMI

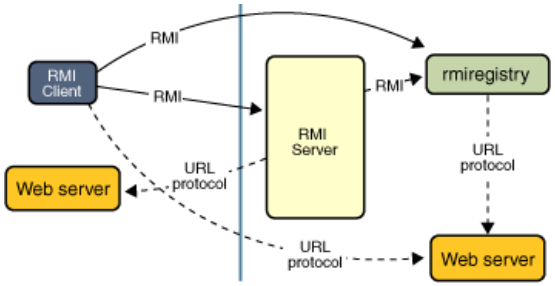

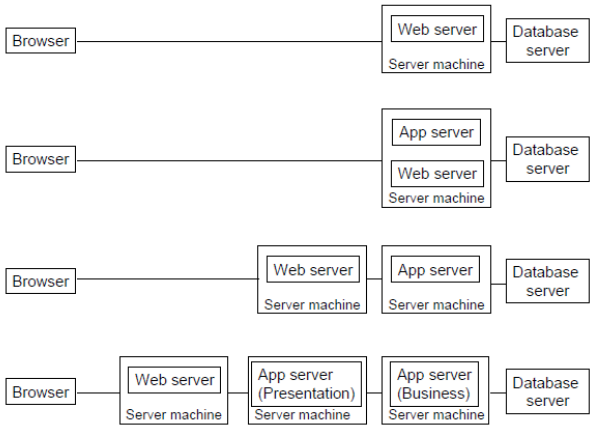

Evolution of Web and App Servers

Web Services with REST API

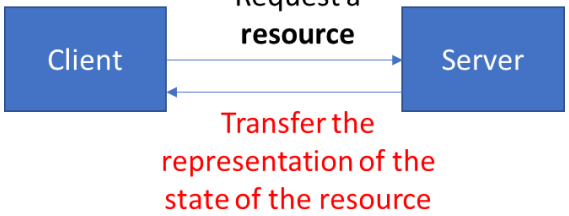

oAuth

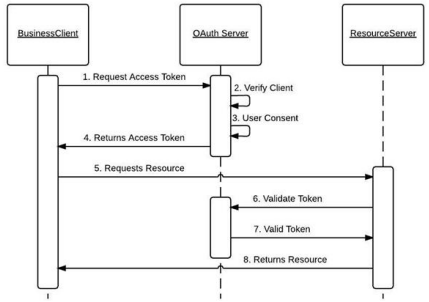

Rate Limiting

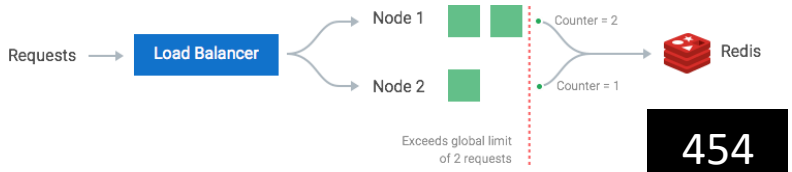

Building Web Services

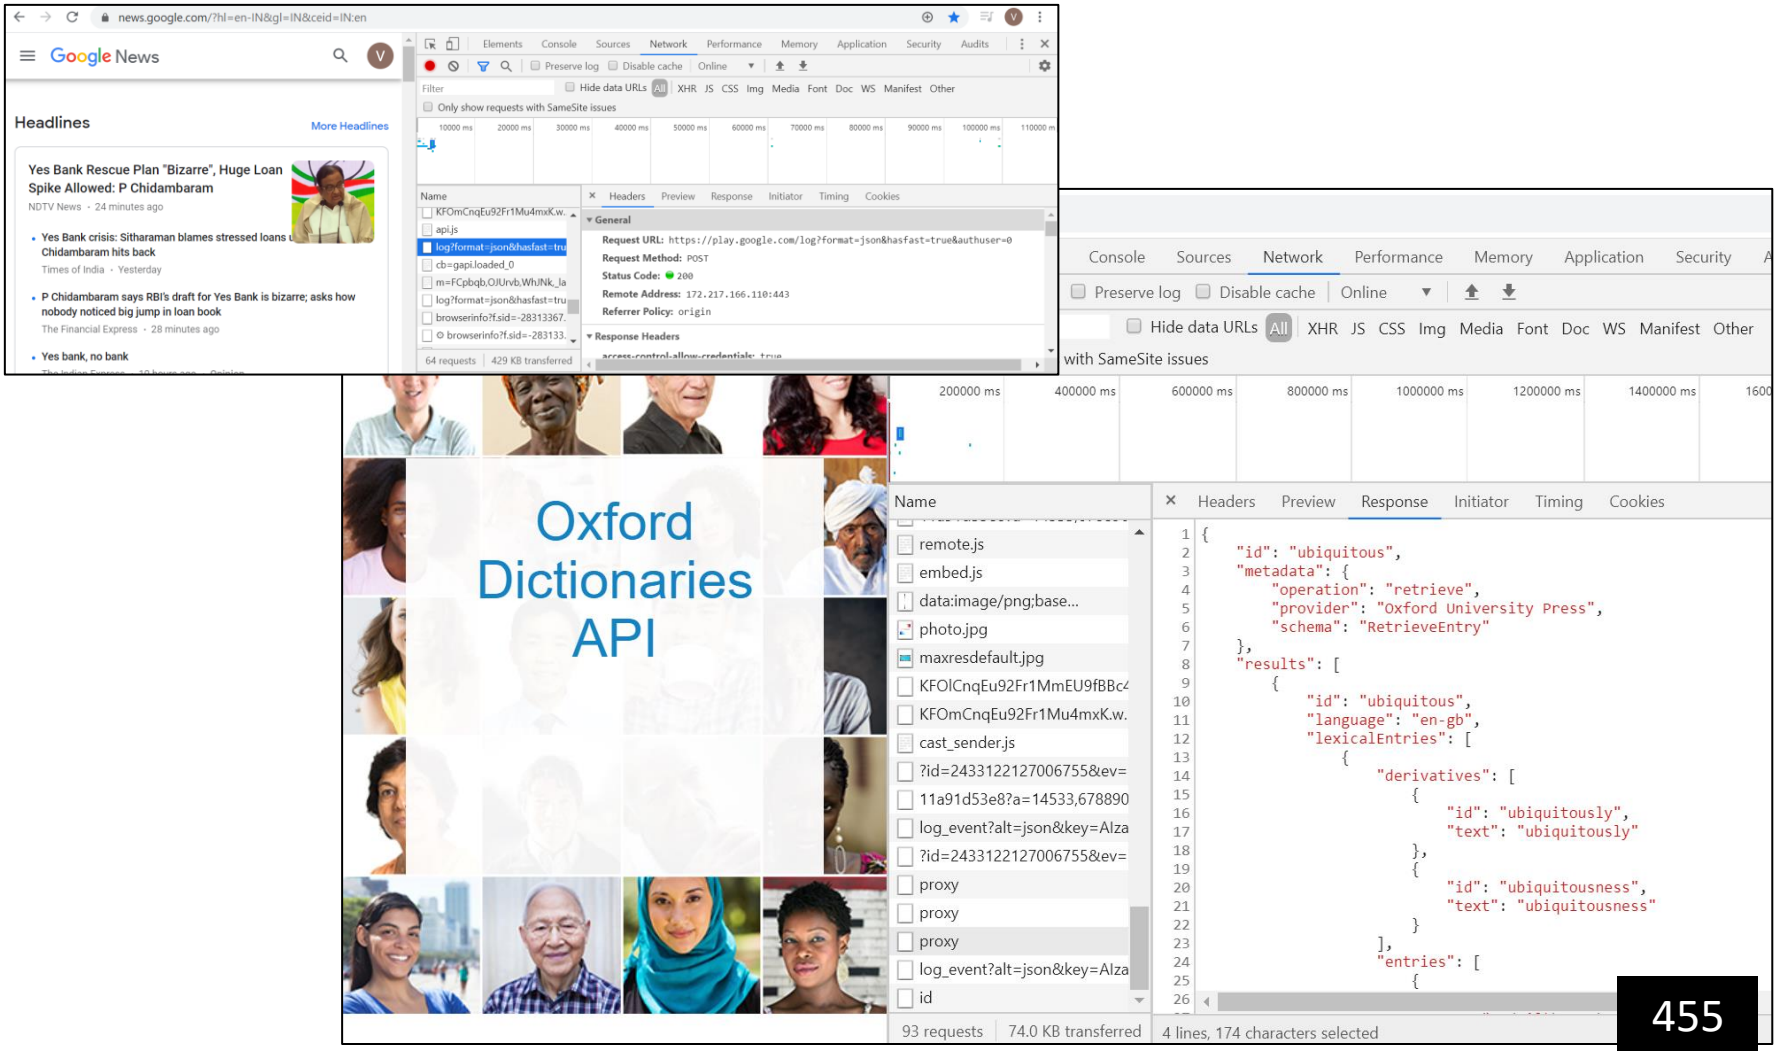

Thank You

Please remember to give elaborate course feedback. I take my course feedback seriously to improve teaching quality including but not limited to the content, presentation materials, and delivery.